

A fork() in the Road

Jeremy Mates

June 24, 2021

In “A fork() in the road” Baumann et al. [2019] two programs are mentioned as being slow with with fork(2)—Chrome, and node. These are atypical choices, as we shall see. Chrome uses abnormal amounts of memory, and a good security policy might well deny Chrome the right to fork if, hypothetically, Chrome were to suffer from some security issue. What should Chrome do if it cannot fork? Perhaps it should, at launch, fork a small supervisor process—maybe called `init`—which would then be responsible for starting the other processes that Chrome needs—perhaps via some “system daemon” framework. This way, processes with heavy memory use or multithreading needs would only do that, and would not ever be calling fork or exec. Did I just describe an operating system? Probably. Is Chrome a not very good operating system? Probably.

node used to crash (in 2019, but no longer does in v12.16.1) if you printed to the console in a loop via

```
$ node -e 'while (1) { process.stdout.write("a") }' > /dev/null
Bus error (core dumped)
```

so it’s not surprising to me that it would have other problems such as taking a long time to fork. Perhaps as skill with heavily multithreaded programs goes up such issues will get ironed out? Assuming the process even needs to be multithreaded and to fork.

Anyways, how atypical is Chrome? Let’s look at memory use on the 2009 MacBook.

PID	COMMAND	%CPU	TIME	#TH	#WQ	#PORT	MEM
0	kernel_task	0.7	12:14.54	114/2	0	2	380M
380	QuickTime Pl	0.0	50:45.86	15	2	342	57M
232	WindowServer	0.3	47:39.50	5	2	311	56M
281	iTerm2	0.2	03:04.65	6	0	252	28M
287	Finder	0.0	00:07.63	3	0	271	26M
285	Dock	0.0	00:01.76	3	0	225	24M
463	Sudoku	0.0	00:09.48	6	1	275	22M
482	Dictionary	0.0	00:00.79	5	0	200	19M
298	Spotlight	0.0	00:00.73	4	0	171	10M

Do those memory intensive processes fork? At all? What does fork a lot? The shells, which are using. . . oh there they are, 1888K, 2268K, etc. Would `posix_spawn` even help for processes that small? How much memory does Chrome use? I’d have to install it to find out; Safari.app under Mac OS X 10.11 uses around 128M doing nothing, or two orders of magnitude more memory than the shells. I’ve heard stories of Chrome eating into another order of magnitude, or two? Chrome having problems with fork() is less about fork() and more about Chrome being, well, bloaty.

Cherry Picking

I too can cherry pick things favorable to my position. Here are my OpenBSD processes that fork, a lot:

```
$ ps axo rss,vsz,args | grep s[h]
968  1008  /bin/sh /etc/X11/xenodm/Xsession
1148  1220  /bin/ksh -l
1052  1020  /bin/ksh -l
1068  1152  /bin/ksh -l
1036  1016  /bin/ksh -l
```

And here we sort on the RES column in `top(1)` with the output cleaned up a bit:

PID	USERNAME	SIZE	RES	STATE	TIME	CPU	COMMAND
64699	_x11	42M	52M	sleep/0	0:19	0.49%	Xorg
54110	jmates	19M	24M	sleep/0	0:02	0.00%	mupdf-x11
49777	jmates	12M	18M	sleep/0	0:04	0.00%	irssi
2134	root	2100K	8080K	idle	0:01	0.00%	xenodm
3606	jmates	1148K	6748K	sleep/1	0:00	0.00%	cwm
45572	jmates	1884K	5784K	sleep/1	0:05	0.00%	xterm
14369	jmates	1884K	5756K	sleep/1	0:07	0.00%	xterm
86394	jmates	1868K	5740K	sleep/1	0:01	0.00%	xterm
46500	jmates	1864K	5740K	sleep/0	0:01	0.00%	xterm
76013	jmates	4812K	5344K	sleep/1	0:01	0.00%	vim

`cwm(1)` and `vim(1)` are the only where notable fork/exec would happen, and not ever in any sort of performance critical path; in the others fork/exec would be minimal to zero; in the case of Xorg and `irssi` `execve(2)` is not even allowed. And the multi-threaded programs, all of them:

```
49777 irssi
64699 /usr/X11R6/bin/X
```

The Replacements

Assuming, hypothetically, that we are on some sort of unix that runs memory-intensive processes with heavy multithreading that also need to fork, what are we to replace fork with? The paper indicates `posix_spawn`, which is “not a complete replacement for fork and exec... it... lacks an effective error reporting mechanism”[Baumann et al., 2019, p.5]. Perhaps someone could spend time (isn’t programmer time expensive?) to catalog what existing uses of fork could actually be rewritten—and then debugged, and tests written for any regressions (wait, are there even tests?), the documentation updated, etc.—to use `posix_spawn` and then process startup might be a bit faster? Is this even a problem? `vim(1)` is slow to come up cold off of disk, but that’s mostly due to the spinning metal hard drive and the 0.8 GHz CPU running OpenBSD, an OS well noted for its speed. Here are some hot cache forks and execs:

```
$ repeat 3 time vim -c quit
0.11 real          0.08 user          0.03 sys
0.11 real          0.09 user          0.02 sys
0.11 real          0.08 user          0.03 sys
$ highcpu
$ repeat 3 time vim -c quit
0.04 real          0.04 user          0.01 sys
0.04 real          0.03 user          0.01 sys
0.05 real          0.04 user          0.01 sys
```

How would `posix_spawn` help here? A migration is a non-starter if you need that error reporting, or otherwise looks to be of moderate effort and low reward. Perhaps if `vim(1)` were being spawned often from software with excessive memory use? But why run such? Why not launch `vim(1)` from a small, fast shell?

An Aside with `nmap`

I did once have a problem with launching many `nmap(1)` processes; I was testing a LDAP server, where lots of connections (or traffic, or something) would (sometimes) wedge the server—obviously only in production, and never in the test environment. The test script probably looked something like:

```
for i in ...; do nmap ... & done
```

Except it was much less pretty, maybe with a wait and a second loop to keep spawning more `nmaps`. Did I rewrite this slow thing to use `posix_spawn`? Nope. I wrote a new script—two actually—that performed a test inside a very tight loop in a single process. The shell forked and execed these test programs, and LDAP fell over. Success! Where would `posix_spawn` help? Maybe to shave off some script startup time? Noise,

compared to converting a quick “I have an idea. . .” mess of a shell pipeline (the prototype) into a pair of very fast test scripts (a first implementation).

Back to the so-called Alternatives

`vfork(2)` is mentioned as an alternative BSD thing. Let’s read the fine manual.

BUGS

This system call will be eliminated when proper system sharing mechanisms are implemented. Users should not depend on the memory sharing semantics of `vfork` as it will, in that case, be made synonymous to `fork`.

Based on that I’d agree with the article that “in most cases it is better avoided”[Baumann et al., 2019, p.5]. The documentation on Mac OS X 10.11 is not known for being updated or up-to-date; let’s try OpenBSD.

DESCRIPTION

`vfork()` was originally used to create new processes without fully copying the address space of the old process, which is horrendously inefficient in a paged environment. It was useful when the purpose of `fork(2)` would have been to create a new system context for an `execve(2)`. Since `fork(2)` is now efficient, even in the above case, the need for `vfork()` has diminished.

...

HISTORY

The `vfork()` function call appeared in 3.0BSD with the additional semantics that the child process ran in the memory of the parent until it called `execve(2)` or exited. That sharing of memory was removed in 4.4BSD, leaving just the semantics of blocking the parent until the child calls `execve(2)` or exits.

Not seeing anything new nor viable about `vfork(2)`. Why is it even listed as an alternative?

The article then mentions some low-level something from elsewhere, and concludes that it “seems at first glance challenging, but may also be productive for future research”[Baumann et al., 2019, p.6]. Good luck with that?

How about `clone()`?

% **man clone**

No manual entry for clone

\$ **man clone**

man: No entry for clone in the manual.

So not portable, plus “clone suffers most of the same problems as fork”[Baumann et al., 2019, p.6].

To conclude, there are simply not any viable alternatives; the call to “strongly discourage” the use of `fork` in new code is both laughable and absurd. Why not simply use `fork()` in a small and fast single threaded processes. . . oh wait, the authors did realize that. In hindsight.

References

Andrew Baumann, Jonathan Appavoo, Orran Krieger, and Timothy Roscoe. Hotos 2019: Proceedings of the workshop on hot topics in operating systems. <https://dl.acm.org/doi/10.1145/3317550.3321435>, 2019.